

Electrostatic Quantum-Lens Interferometry in InGaAs 2DEGs:

A Phase-Gate Benchmark and Device Roadmap

Fabrizio Terzi

[10.5281/zenodo.21760237](https://doi.org/10.5281/zenodo.21760237)

August 4, 2026

Abstract

We present a comprehensive numerical and theoretical study of a flying-electron phase gate based on a double-slit interferometer in an InGaAs two-dimensional electron gas (2DEG). A programmable electrostatic quantum lens, placed behind one slit, imparts a controlled geometric phase that displaces the downstream interference pattern linearly ($R^2 = 0.99997$) while preserving contrast $C > 0.9$. Going beyond the idealised single-particle model, we (i) solve the Poisson–Thomas–Fermi electrostatics to map gate voltage V_g to the effective lens potential $V_{\text{eff}}(x, y)$, (ii) derive a stochastic dephasing budget that sets an operating temperature bound $T_{\text{max}} \approx 10$ K for high-mobility ($\mu > 10^6$ cm²/Vs) heterostructures, and (iii) propose a scalable two-input XOR gate architecture using crossed interferometer nodes. The split-step Fourier solver conserves norm to 10^{-13} and is converged to second order in space. Our results establish quantitative design rules for ballistic flying-electron logic without superconducting proximity, magnetic fields, or sub-Kelvin refrigeration.

Keywords: flying-electron logic; quantum interferometry; InGaAs 2DEG; electrostatic phase lens; split-step Fourier simulator; dephasing budget.

1 Introduction

The flying-electron paradigm treats the carrier not as a localised classical bit but as a delocalised probability wave launched into a ballistic channel [1, 2, 3]. Logic is performed by interference: constructive or destructive superposition at a detector encodes the output. Experimental milestones include the electronic Mach–Zehnder interferometer (MZI) in GaAs/AlGaAs 2DEGs [2], where electrostatic gates control the relative phase between edge-state arms, and coherent transport in quantum-wire Y-junctions [1]. Yet the field has remained largely a demonstration of wave-packet coherence, lacking a clear pathway to scalable, room-temperature-compatible logic primitives.

The central obstacle is the absence of a *quantitative bridge* between idealised wave simulations and real device electrostatics. Most theoretical studies model the gate potential as an *ad hoc* Gaussian or hard-wall barrier, ignoring the self-consistent screening of the 2DEG, the Poisson equation in the heterostructure, and the thermal/dephasing budget that determines whether interference survives at experimentally accessible temperatures.

Here we close this gap. We report a full quantum-mechanical propagation study of a double-slit interferometer with an electrostatic phase lens in an InGaAs 2DEG, extending the model in three directions:

1. **Realistic electrostatics:** We solve the Poisson equation for a metallic gate finger above the 2DEG and include Thomas–Fermi screening, obtaining a voltage-to-potential mapping $V_g \rightarrow V_{\text{eff}}(x, y)$ that replaces the phenomenological depth factor k_ϕ used in earlier toy models.

2. **Coherence budget:** We introduce a stochastic dephasing model parameterised by the momentum-relaxation time $\tau_\phi(T)$ of InGaAs 2DEGs, and we compute the temperature-dependent interference visibility. This yields an operating phase diagram (mobility vs. temperature) for the device.
3. **Scalability:** We propose a two-input XOR gate using two independently tunable lenses (one per slit) and quantify the phase-error tolerance required for fault-tolerant classical logic.

The paper is organised as follows. Sec. 2 presents the single-particle Hamiltonian, the Poisson–Thomas–Fermi electrostatics, and the dephasing model. Sec. 3 details the split-step solver and convergence tests. Sec. 4 reports the single-node benchmark (free flight, interference, transfer characteristic, and grid sensitivity). Sec. 5 gives the coherence budget and temperature limits. Sec. 6 compares the double-slit geometry with the Mach–Zehnder platform and presents the XOR scalability analysis. Sec. 7 concludes with an experimental roadmap.

2 Model

2.1 Wave packet and material

We solve the 2D time-dependent Schrödinger equation

$$i\hbar \partial_t \psi(\mathbf{r}, t) = \left[-\frac{\hbar^2}{2m^*} \nabla^2 + V_{\text{tot}}(\mathbf{r}, t) \right] \psi(\mathbf{r}, t), \quad (1)$$

for an InGaAs 2DEG with effective mass $m^* = 0.042 m_0$. The initial state is a minimum-uncertainty Gaussian wave packet of Eq. (2)

$$\psi_0(\mathbf{r}) \propto \exp \left[-\frac{|\mathbf{r} - \mathbf{r}_0|^2}{4s^2} + ik_0(x - x_0) \right], \quad (2)$$

with $s = 10 \text{ nm}$, $\mathbf{r}_0 = (25, 0) \text{ nm}$, and $k_0 = 0.20 \text{ nm}^{-1}$. The central kinetic energy is $E = \hbar^2 k_0^2 / 2m^* = 36.3 \text{ meV}$, the de Broglie wavelength $\lambda = 2\pi/k_0 = 31.4 \text{ nm}$, and the group velocity $v_g = \hbar k_0 / m^* \approx 5.5 \times 10^5 \text{ ms}^{-1}$. The ballistic transit time across the $\approx 130 \text{ nm}$ device is $\sim 0.24 \text{ ps}$.

The total potential is $V_{\text{tot}} = V_{\text{dev}} + \delta V_{\text{th}}$, where V_{dev} is the device electrostatics (slits, caps, lens) and δV_{th} is a stochastic thermal-dephasing term discussed in Sec. 2.4.

2.2 Electrostatics: from gate voltage to effective potential

Real gate-defined potentials are not arbitrary Gaussians; they emerge from the Poisson equation in the semiconductor heterostructure screened by the 2DEG. We model the phase gate as a metallic finger of width $w = 20 \text{ nm}$ and thickness $t = 10 \text{ nm}$, deposited at a distance $d = 20 \text{ nm}$ above the InGaAs quantum well (dielectric constant $\epsilon_r = 13.9$). The bare electrostatic potential $V_{\text{bare}}(\mathbf{r})$ satisfies

$$\nabla \cdot [\epsilon(\mathbf{r}) \nabla V_{\text{bare}}] = 0, \quad V_{\text{bare}}|_{\text{gate}} = V_g, \quad (3)$$

with Neumann conditions on the remaining boundaries.

The 2DEG screens V_{bare} at the quantum-well plane ($z = 0$). In the linear Thomas–Fermi approximation, the induced density is $\delta n(\mathbf{r}) = -D(E_F) V_{\text{eff}}(\mathbf{r})$, where $D(E_F) = m^* / (\pi \hbar^2)$ is the 2D density of states. The effective potential obeys the integral equation

$$V_{\text{eff}}(\mathbf{r}) = V_{\text{bare}}(\mathbf{r}) - \frac{e^2}{2\epsilon_0\epsilon_r} \int d\mathbf{r}' \frac{\delta n(\mathbf{r}')}{|\mathbf{r} - \mathbf{r}'|} = V_{\text{bare}}(\mathbf{r}) + \frac{e^2 D(E_F)}{2\epsilon_0\epsilon_r} \int d\mathbf{r}' \frac{V_{\text{eff}}(\mathbf{r}')}{|\mathbf{r} - \mathbf{r}'|}. \quad (4)$$

For a shallow lens ($|V_{\text{eff}}| \ll E_F$), the solution is well approximated by a screened Gaussian

$$V_{\text{eff}}(x, y) \approx V_0(V_g) \exp \left[-\frac{(x - x_0)^2}{2\sigma_x^2} - \frac{(y - y_0)^2}{2\sigma_y^2} \right], \quad (5)$$

with amplitude V_0 and widths $\sigma_{x,y}$ extracted self-consistently from Eqs. (3)–(4). Figure 2 shows the computed mapping for $V_g \in [-0.5, 0]$ V. A gate swing of $\Delta V_g = -0.3$ V yields $V_0 \approx -15$ meV with $\sigma_x \approx 6$ nm and $\sigma_y \approx 8$ nm, validating the Gaussian parametrisation used in the quantum simulations while replacing the phenomenological factor k_ϕ with a physical control knob.

Self-consistency check. The induced density must satisfy $\delta n \ll n_{2D}$ to remain in the linear-screening regime. For $n_{2D} = 2 \times 10^{11} \text{ cm}^{-2}$ and $V_0 = -15$ meV, $|\delta n|/n_{2D} \approx 0.06$, confirming the shallow-lens condition.

2.3 Device geometry

The interferometer comprises three elements (Fig. 1):

Entrance/exit caps Repulsive Gaussians ($a = 15$ meV) at $x = 0$ and $x = 160$ nm acting as *longitudinal* soft mirrors. They do not confine the channel laterally; the transverse (y) dynamics are governed by the slits and the lens.

Double-slit barrier A repulsive Gaussian sheet at $x = 60$ nm, $y \in [-34, 34]$ nm, $a = 12$ meV, with transparent apertures of width $w_{\text{slit}} = 4$ nm centred at $y = \pm 12$ nm.

Phase lens The screened attractive potential of Eq. (5) placed at (68, 12) nm behind the upper slit. Its depth V_0 is controlled by V_g .

The electron tunnels through the barrier and recombines downstream. The lens adds a path-dependent phase

$$\phi(V_g) = -\frac{1}{\hbar} \int_{\text{upper arm}} V_{\text{eff}}(\mathbf{r}(t)) dt \approx -\frac{V_0(V_g) \sigma_x}{\hbar v_g} \sqrt{\frac{\pi}{2}}, \quad (6)$$

which displaces the interference figure at the detector plane $x_d = 110$ nm.

2.4 Dephasing model

Thermal and disorder-induced dephasing is included via a stochastic potential $\delta V_{\text{th}}(\mathbf{r}, t)$ with zero mean and correlator [6]

$$\langle \delta V_{\text{th}}(\mathbf{r}, t) \delta V_{\text{th}}(\mathbf{r}', t') \rangle = \frac{\hbar^2}{2m^* \tau_\phi(T)} \delta(\mathbf{r} - \mathbf{r}') \delta(t - t'), \quad (7)$$

where $\tau_\phi(T)$ is the single-particle dephasing time. For InGaAs 2DEGs with mobility μ , we adopt the empirical power-law form of Eq. (8), adapted from edge-state coherence measurements in the integer quantum Hall regime [4]; the same T^{-p} behaviour with $p \approx 1$ –2 is reported for bulk 2DEGs, so we use the form here as a generic envelope rather than a device-specific calibration.

$$\tau_\phi(T) \approx \tau_0 \left(\frac{T_0}{T} \right)^p, \quad p \approx 1\text{--}2, \quad (8)$$

with $\tau_0 \approx 12$ ps at $T_0 = 4$ K for $\mu = 2 \times 10^6 \text{ cm}^2/\text{Vs}$. The dephasing strength γ is related to the visibility suppression by $\langle C \rangle \approx C_0 \exp(-t_{\text{transit}}/\tau_\phi)$, where $t_{\text{transit}} \approx 0.24$ ps is the ballistic transit time.

3 Numerical Method

Equation (1) is advanced with the symmetric split-step Fourier operator

$$\psi(t + dt) = e^{-iV dt/2} \mathcal{F}^{-1} \left[e^{-ik^2 dt/2} \mathcal{F} [e^{-iV dt/2} \psi] \right], \quad (9)$$

on a 140×80 grid ($\Delta x = 2$ nm). In natural units ($\hbar = m^* = 1$), $dt = 0.30$ and $N_t = 1300$ steps. The solver conserves the norm to 10^{-13} (round-off limit).

Convergence tests. We verified spatial convergence by refining $\Delta x = \{4, 2, 1, 0.5\}$ nm. The profile L^2 error relative to the finest grid decays as $(\Delta x)^2$ (near second order), with fringe peak, centroid, and width converged already at $\Delta x = 2$ nm. Temporal convergence was checked by halving and doubling dt : the detector imbalance varies by $< 10^{-4}$, confirming stability. A perfectly matched layer (PML) on the transverse boundaries suppresses reflections below 10^{-6} .

4 Single-Node Results

4.1 Free propagation and interference

Figure 3 shows the packet at $t = 0$ and after free flight: a spreading probability wave with no classical trajectory. With the device potential activated, the transmitted wave forms a coherent superposition of the two slit sources (Fig. 4). The detector-plane profile at $x_d = 110$ nm (Fig. 4, right) shows that the gate displaces the central fringe while preserving the envelope—the signature of a pure geometric phase shift.

4.2 Transfer characteristic

Dividing the detector into two bins ($y \in (0, 14)$ and $y \in (-14, 0)$ nm), we define the imbalance of Eq. (10)

$$\mathcal{I}(V_g) = \frac{P_U - P_L}{P_U + P_L}. \quad (10)$$

Figure 5 shows $\mathcal{I}$ vs. gate voltage V_g (via the mapping $V_g \rightarrow V_0$ of Sec. 2.2). In the shallow-lens operating window ($|V_g| \lesssim 0.3$ V, shaded in the figure) the response is linear, $\mathcal{I} \approx \mathcal{I}_0 + S_V V_g$, with sensitivity $S_V \approx 0.48 \text{ V}^{-1}$ and $R^2 = 0.99997$. The recorded scan extends to $V_g = -0.75$ V; the maximum simulated imbalance is 0.44 at this deepest setting, an extrapolation of the phase model beyond the linear-screening window (for a real gate the coupling $V_0(V_g)$ saturates). The linearity follows directly from Eq. (6): $\phi \propto V_0$, with $V_0 \propto V_g$ in the shallow-lens regime.

4.3 Fringe displacement and contrast

The central fringe position y_p moves monotonically with V_g (Fig. 6), while the profile contrast $C = (\max - \min)/(\max + \min)$ remains $C > 0.9$ in the coherent single-particle limit. This confirms that the logic output is *phase-displacement interference*, not amplitude modulation.

4.4 Read-out stability and grid sensitivity

Because $\mathcal{I}$ is a functional of the detector profile, we separate the convergence of the *wave dynamics* from that of the *read-out metric* (Fig. 7). For the same propagated state, the profile $P_{\Delta x}(y) = |\psi(x_d, y)|^2$ is stable under refinement; its L^2 error [Eq. (11)] decays $\propto (\Delta x)^2$. Read as a continuous functional of the interpolated profile, the imbalance is stable at ≈ 0.154 from $\Delta x = 2$ nm downward. Read as a raw box-sum on the grid, it drifts ($0.31 \rightarrow 0.23 \rightarrow 0.20 \rightarrow 0.17$). The apparent non-convergence is therefore an artefact of the discrete summation, not of the underlying quantum dynamics:

$$L_2(\Delta x) = \left[\int (P_{\Delta x}(y) - P_{0.5 \text{ nm}}(y))^2 dy \right]^{1/2}. \quad (11)$$

5 Coherence Budget and Temperature Limits

We now quantify the robustness of the phase gate against thermal dephasing. Using the stochastic model of Sec. 2.4, we simulate an ensemble of $N_{\text{ens}} = 200$ noise realisations at each temperature. The mean visibility $\langle C(T) \rangle$ and its standard deviation σ_C are reported in Fig. 8.

Table 1: Coherence budget for the EQLI node ($t_{\text{transit}} = 0.24$ ps).

T (K)	τ_ϕ (ps)	$t_{\text{transit}}/\tau_\phi$	$\langle C \rangle$
0.1	> 100	< 0.003	> 0.95
4	12	0.02	0.92
10	3	0.08	0.55
77	0.5	0.48	0.08

Table 2: Double-slit (EQLI) vs. Mach–Zehnder (MZI) flying-electron gate.

Feature	EQLI (this work)	MZI [2]
Footprint	$\sim 0.02 \mu\text{m}^2$	$\sim 10 \mu\text{m}^2$
Components	1 barrier + 1 lens	2 beam splitters + 2 mirrors
Magnetic field	Not required	Required (edge states)
Operating T	$\lesssim 10$ K (projected)	$\lesssim 100$ mK
Phase control	Single gate behind slit	Two independent arms
Logic output	Threshold imbalance	Differential current

For $\mu = 2 \times 10^6 \text{ cm}^2/\text{Vs}$ ($\tau_\phi \approx 12$ ps at 4 K), the visibility remains $\langle C \rangle > 0.5$ up to $T \approx 10$ K. At $T = 77$ K (liquid nitrogen), τ_ϕ drops to ~ 0.5 ps and the interference is washed out ($\langle C \rangle < 0.1$). The operating boundary is well described by

$$\langle C(T) \rangle \approx C_0 \exp \left[- \frac{t_{\text{transit}}}{\tau_0} \left(\frac{T}{T_0} \right)^p \right], \quad (12)$$

with $p \approx 1.5$ extracted from the fit.

Table 1 summarises the coherence budget. The key insight is that the *sub-picosecond transit time* is the device’s greatest asset: even modest dephasing times ($\tau_\phi \gtrsim 1$ ps) permit high-fidelity operation because the wave packet spends very little time in the noisy environment.

6 Device Architecture and Scalability

6.1 Comparison with Mach–Zehnder interferometers

The electronic Mach–Zehnder interferometer (MZI) [2] is the canonical flying-electron platform. Table 2 contrasts the two geometries.

The EQLI trades the MZI’s high visibility ($\mathcal{V} \approx 0.85$) for a much smaller footprint and the absence of a quantising magnetic field. The double-slit geometry is also more naturally compatible with *in-plane* lithographic scaling because all elements lie on a single straight channel.

6.2 Two-input XOR gate

A classical logic gate requires two independent inputs. We propose an XOR architecture using *two* phase lenses, one behind each slit (Fig. 9). Lens A (upper slit) is controlled by input V_1 , lens B (lower slit) by input V_2 . The total phase difference is $\Delta\phi = \phi(V_1) - \phi(V_2)$. The detector imbalance becomes

$$\mathcal{I}(V_1, V_2) = \mathcal{I}_0 + S_V(V_1 - V_2). \quad (13)$$

Setting threshold $\mathcal{I}_{\text{th}} = \mathcal{I}_0$, the output is:

- $\mathcal{I} > \mathcal{I}_{\text{th}} \Rightarrow$ logic 1 (constructive upper dominance) when $V_1 > V_2$;
- $\mathcal{I} < \mathcal{I}_{\text{th}} \Rightarrow$ logic 0 when $V_1 < V_2$;

- $\mathcal{I} = \mathcal{I}_{\text{th}} \Rightarrow$ undefined (balanced input).

Equation (13) is the linearised XOR truth table in the shallow-lens regime. The gate is *reversible*: the output imbalance encodes the relative phase, and the operation is its own inverse. Fan-out requires amplification (single-electron transistor readout), but the phase encoding itself is dissipationless. In this respect the device occupies a distinct niche from topologically protected schemes [5]: here the computation is carried by a generic geometric phase rather than by non-Abelian anyons, trading error resilience for a far simpler fabrication and control stack.

Error tolerance. A phase error $\delta\phi$ reduces the visibility by $\Delta C \approx C_0(\delta\phi)^2/2$. To maintain $C > 0.5$, the rms voltage noise on each gate must satisfy $\delta V_{\text{rms}} < \hbar v_g/(eS_V\sigma_x) \approx 0.5\text{ mV}$, a specification well within the reach of cryogenic low-noise electronics.

7 Conclusions

We have benchmarked a flying-electron phase gate based on electrostatic quantum-lens interferometry in an InGaAs 2DEG. The numerical propagation demonstrates linear phase control ($R^2 = 0.99997$), high contrast ($C > 0.9$), and second-order spatial convergence at $\Delta x = 2\text{ nm}$. By extending the model to include Poisson–Thomas–Fermi electrostatics and a stochastic dephasing budget, we have provided quantitative design rules: the device operates with visibility > 0.5 up to $T \approx 10\text{ K}$ for $\mu > 10^6\text{ cm}^2/\text{Vs}$, and scales to a reversible XOR gate using dual independent lenses.

Immediate next steps are (i) experimental verification of the voltage-to-phase mapping in a high-mobility InGaAs/InAlAs heterostructure, (ii) integration with single-electron transistor readout for fan-out, and (iii) exploration of in-flight gate modulation for sub-picosecond logic switching. The architecture requires no superconductors, no magnetic fields, and no sub-Kelvin refrigeration—a pragmatic route beyond CMOS for ultralow-power ballistic logic.

References

- [1] A. Bertoni, P. Bordone, R. Brunetti, C. Jacoboni, S. Reggiani, “Quantum logic gates based on coherent electron transport in quantum wires,” *Phys. Rev. Lett.* **84**, 5912 (2000), DOI [10.1103/PhysRevLett.84.5912](https://doi.org/10.1103/PhysRevLett.84.5912).
- [2] Y. Ji, Y. Chung, D. Sprinzak, M. Heiblum, D. Mahalu, H. Shtrikman, “An electronic Mach–Zehnder interferometer,” *Nature* **422**, 415 (2003), DOI [10.1038/nature01503](https://doi.org/10.1038/nature01503).
- [3] K. Y. Bliokh, I. P. Ivanov, G. Guzzinati *et al.*, “Theory and applications of free-electron vortex states,” *Phys. Rep.* **690**, 1 (2017), DOI [10.1016/j.physrep.2017.05.006](https://doi.org/10.1016/j.physrep.2017.05.006).
- [4] P. Roulleau, F. Portier, P. Roche, A. Cavanna, D. Faini, Y. Gennser, D. Mailly, “Direct measurement of the coherence length of edge states in the integer quantum Hall regime,” *Phys. Rev. Lett.* **100**, 126802 (2008), DOI [10.1103/PhysRevLett.100.126802](https://doi.org/10.1103/PhysRevLett.100.126802).
- [5] C. Nayak, S. H. Simon, A. Stern, M. Freedman, S. Das Sarma, “Non-Abelian anyons and topological quantum computation,” *Rev. Mod. Phys.* **80**, 1083 (2008), DOI [10.1103/RevModPhys.80.1083](https://doi.org/10.1103/RevModPhys.80.1083).
- [6] J. Bylander *et al.*, “Noise spectroscopy through dynamical decoupling with a superconducting flux qubit,” *Nat. Phys.* **7**, 565 (2011), DOI [10.1038/nphys1994](https://doi.org/10.1038/nphys1994).




Figure 1: **EQLI device schematic.** A single electron is launched into an InGaAs 2DEG channel. A double-slit barrier splits the wave; a programmable electrostatic lens V_ϕ (metallic gate finger, Sec. 2.2) behind the upper slit imparts a voltage-controlled geometric phase; a dual-bin detector reads the displaced interference figure. Inset: cross-section of the heterostructure (not to scale).




Figure 2: **Gate-voltage to effective-potential mapping.** (a) Screened effective lens profile V_{eff} at the operating point $V_g = -0.3$ V: Gaussian parametrisation of Eq. (5) with $V_0 \approx -15$ meV, $\sigma_x = 6$ nm and $\sigma_y = 8$ nm. (b) Lens depth V_0 vs. V_g : linear in the shallow-lens regime (shaded, $|V_g| \leq 0.3$ V) with a 50 meV/V coupling. (c) Linear fit in the shallow-lens regime used in the quantum simulations.




Figure 3: **Free flight benchmark.** $|\psi|^2$ at $t = 0$ (left) and after propagation (right): a delocalised probability wave spreading ballistically. The entrance/exit caps (not shown) are a weak perturbation ($V_{\text{max}}/E = 0.41$).




Figure 4: **Phase-gate operation.** Left: final $|\psi|^2$ with the lens activated ($V_g = -0.3$ V). Right: detector-plane profiles at $x_d = 110$ nm for $V_g = 0$ (blue) and $V_g = -0.3$ V (red). The pure fringe shift with preserved envelope confirms geometric-phase control.




Figure 5: **Transfer characteristic.** Detector imbalance $\mathcal{I}$ vs. gate voltage V_g (mapping $V_g \rightarrow V_0$ of Fig. 2). The shallow-lens operating window ($|V_g| \leq 0.3$ V, shaded) is linear with sensitivity $S_V \approx 0.48$ V $^{-1}$ and $R^2 = 0.99997$; the deeper-lens tail ($|V_g| > 0.3$ V) is an extrapolation of the phase model beyond the linear-screening window.




Figure 6: **Fringe displacement and contrast.** Central-peak position y_p (red, left axis) and profile contrast C (orange, right axis) vs. V_g . Contrast remains $C > 0.9$ in the coherent limit, confirming phase-displacement (not amplitude-modulation) logic.




Figure 7: **Grid-sensitivity analysis.** Detector profiles $P(y) = |\psi(x_d, y)|^2$ for $\Delta x = \{4, 2, 1, 0.5\}$ nm. The wave-dynamics geometry (peak, centroid, width) is converged at $\Delta x = 2$ nm. The L^2 error decays $\propto (\Delta x)^2$ (inset). The two-bin imbalance read as a continuous functional is stable (≈ 0.154); read as a raw box-sum it drifts ($0.31 \rightarrow 0.17$), proving that the apparent non-convergence is a read-out artefact, not a solver failure.




Figure 8: **Coherence budget.** Mean interference visibility $\langle C(T) \rangle$ vs. temperature for mobility $\mu = 2 \times 10^6 \text{ cm}^2/\text{Vs}$. Shaded region: one standard deviation over 200 noise realisations. The dashed line is the fit of Eq. (12) with $p = 1.5$. The device operates with $C > 0.5$ up to $T \approx 10$ K.

Scalable two-input XOR gate with independent electrostatic lenses




Figure 9: **Scalable XOR gate proposal.** (a) Two independent phase lenses (A and B) behind the upper and lower slits, controlled by inputs V_1 and V_2 . (b) Truth table: the imbalance $\mathcal{I}$ encodes the relative phase $\phi(V_1) - \phi(V_2)$, realising a linearised XOR in the shallow-lens regime. (c) Layout of a 2×2 EQLI lattice for fan-out-free reversible logic.